Evolution of Artificial Intelligence From Symbolic Thinking to Agentic AI

The Odyssey of Intelligence: A History of AI

Exploring the evolution of Artificial Intelligence from its conceptual origins to modern Agentic AI.

Highlighting key milestones, paradigm shifts, and technological breakthroughs.

Understanding the foundational principles and the mechanics behind Large Language Models (LLMs).

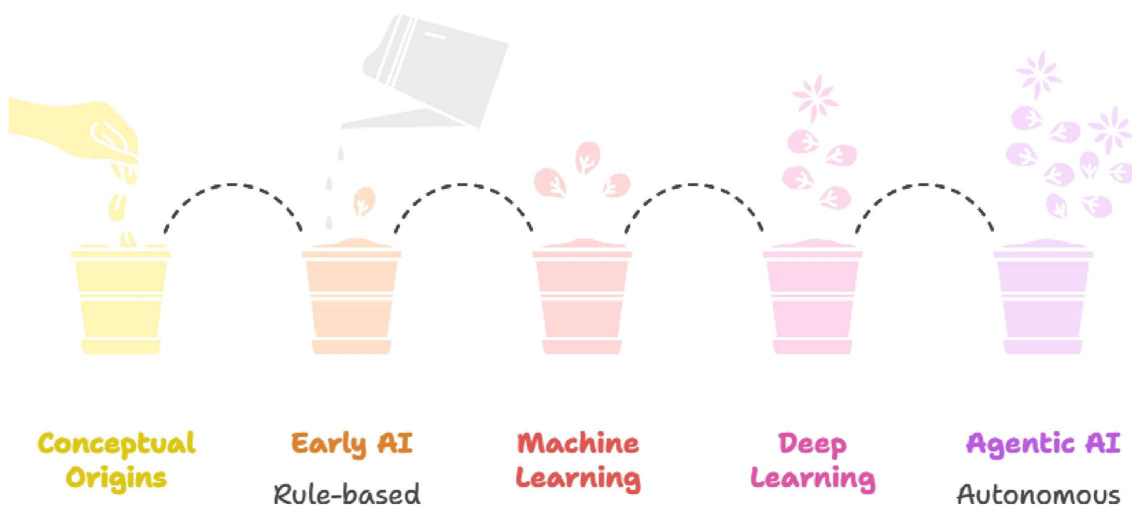

Speaker Notes

Good morning/afternoon everyone. Today, we embark on a fascinating journey through the history of Artificial Intelligence. We'll trace its path from early philosophical dreams to the sophisticated agentic systems we see emerging today. Our discussion will cover the major turning points, the underlying technologies, and a deep dive into how Large Language Models, a truly transformative technology, actually work.

Ancient Dreams and Philosophical Roots

`Slide 2`

Early myths and legends featuring automatons and intelligent constructs (e.g., Golem, Talos).

Philosophical inquiries into the nature of thought, knowledge, and reasoning (e.g., Aristotle's Syllogisms, Descartes' mind-body problem).

Mathematical logic and calculable reasoning laid the groundwork (e.g., Leibniz's 'calculus ratiocinator').

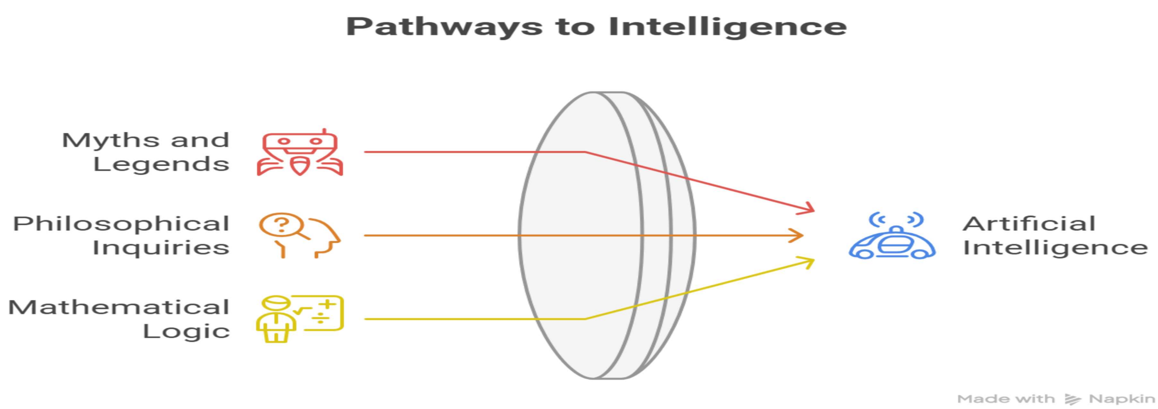

Speaker Notes

The idea of creating intelligent machines is not new; it's a concept deeply embedded in human history, appearing in ancient myths and philosophical texts. From the Golem of Jewish folklore to the bronze giant Talos in Greek mythology, humans have long dreamt of constructing beings that mimic or surpass human capabilities. Philosophers like Aristotle and Leibniz began to formalize reasoning itself, laying a crucial abstract foundation for what would later become computational logic.

The Dawn of AI: Foundational Concepts (1940s-1950s)

`Slide 3`

McCulloch-Pitts artificial neurons (1943): First mathematical model of a neuron.

Norbert Wiener's Cybernetics (1948): The study of control and communication in animals and machines.

Alan Turing's 'Computing Machinery and Intelligence' (1950): Introduced the 'Imitation Game' (Turing Test).

Dartmouth Workshop (1956): Coined the term 'Artificial Intelligence' and marked its formal birth as a field.

Speaker Notes

The mid-20th century saw the pivotal transition from philosophical musings to concrete scientific endeavors. McCulloch and Pitts gave us the first mathematical model of a neuron, suggesting how simple logical operations could be performed. Norbert Wiener's work on Cybernetics provided a framework for understanding self-regulating systems. But it was Alan Turing, with his seminal paper and the Turing Test, who truly set the stage for evaluating machine intelligence. The Dartmouth Workshop in 1956, bringing together pioneers like John McCarthy, Marvin Minsky, and Claude Shannon, formally established AI as a distinct field of study.

The Golden Years & Early Enthusiasm (1950s-1960s)

Slide 4

Logic-based AI and symbolic reasoning dominated initial research.

First AI programs: Logic Theorist (1956) and General Problem Solver (GPS) (1957) by Newell, Simon, and Shaw.

Development of LISP (1958) by John McCarthy, a foundational programming language for AI.

Early successes in chess-playing programs and natural language processing (e.g., ELIZA).

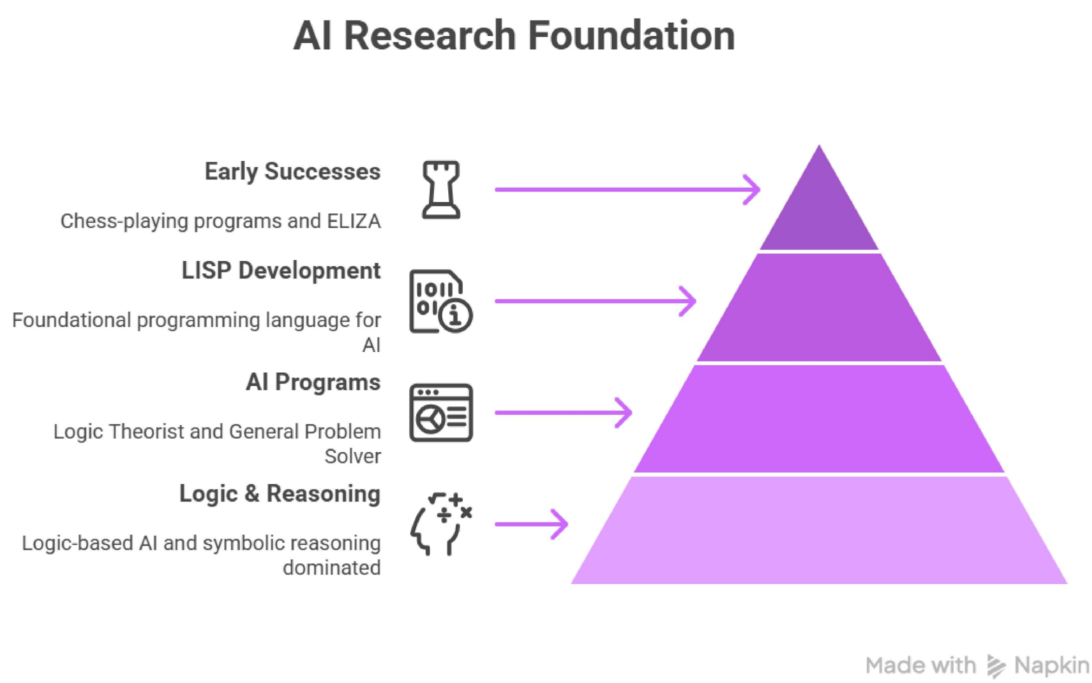

Speaker Notes

Following Dartmouth, the field experienced a period of immense optimism, often referred to as the 'Golden Years.' Researchers believed that human-level intelligence was just around the corner, primarily through symbolic AI. Programs like Logic Theorist and GPS demonstrated that computers could perform complex reasoning tasks. LISP became the language of choice for AI researchers, enabling the manipulation of symbols. Early programs like ELIZA, though simple, captivated the public with their ability to engage in rudimentary conversations, fueling the excitement.

The First AI Winter (1970s)

`Slide 5`

Realization of limitations: AI systems struggled with common sense and scalability.

Computational complexity of symbolic methods became apparent.

Lighthill Report (1973) in the UK highlighted lack of progress, leading to significant funding cuts.

Early promises of 'general intelligence' proved overly ambitious and unattainable with current technology.

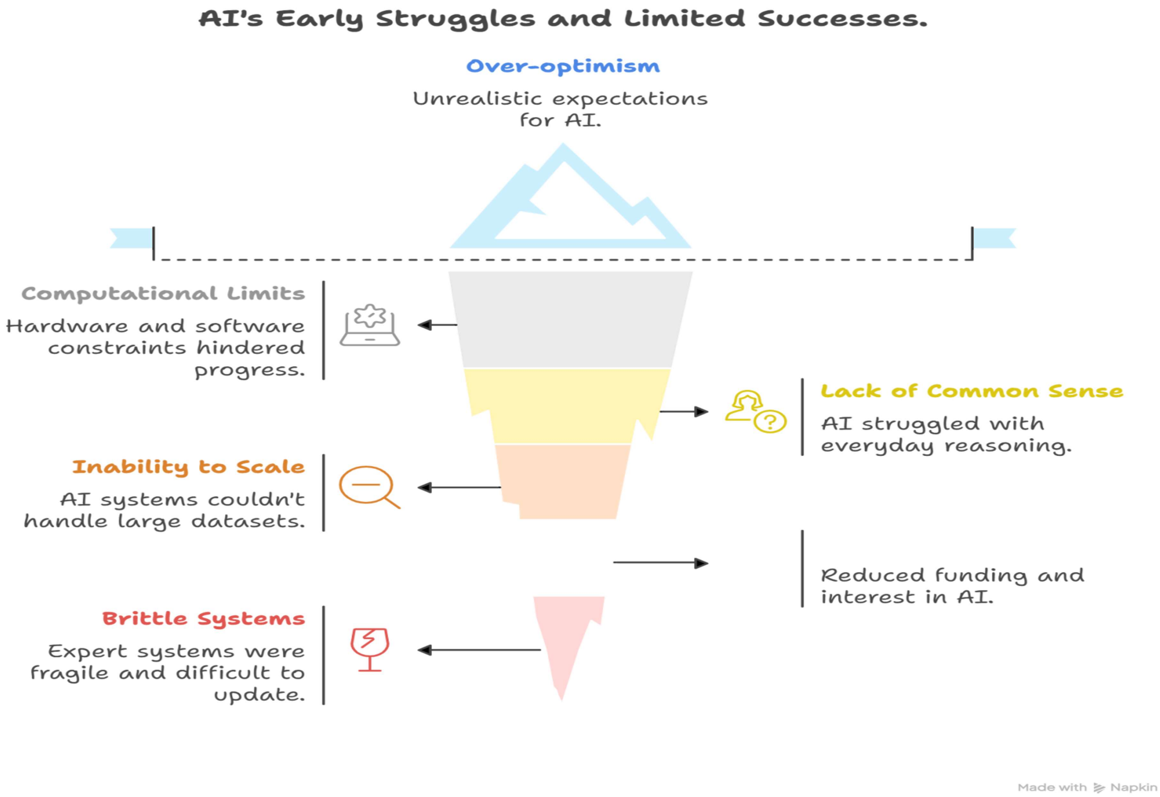

Speaker Notes

However, the initial enthusiasm began to wane as the inherent difficulties of creating truly intelligent systems became clear. AI programs were brittle; they could only operate effectively within narrow, predefined domains and lacked common sense. The computational demands of symbolic reasoning often exceeded the capabilities of the hardware. The Lighthill Report in the UK was a critical turning point, leading to a dramatic reduction in government funding for AI research, marking the beginning of the 'First AI Winter.' This period taught researchers valuable lessons about the challenges of scaling AI beyond toy problems.

The Rise of Expert Systems (1980s)

`Slide 6`

Shift towards 'knowledge-based' systems, focusing on specific domain expertise.

Expert systems like MYCIN (medical diagnosis) and XCON (computer configuration) achieved commercial success.

Relied on explicit knowledge representation through 'if-then' rules and inference engines.

Demonstrated practical value in constrained environments, revitalizing AI research and investment.

Speaker Notes

The 1980s saw a resurgence of AI, driven by the success of expert systems. Instead of aiming for general intelligence, these systems focused on capturing and applying human expertise within *narrow domains. Programs like MYCIN, for diagnosing blood infections, and XCON, for configuring VAX computers, proved highly effective and delivered significant commercial value. This paradigm relied on meticulously hand-coded 'if-then' rules and sophisticated inference engines. While successful, this approach highlighted the immense effort required for knowledge acquisition and the brittleness of these systems outside their specific domains.*

The Second AI Winter & Machine Learning Resurgence (Late 1980s - 1990s)

`Slide 7`

Limitations of expert systems: Difficult to maintain, update, and scale; lacked common sense.

Collapse of the LISP machine market and general disillusionment led to another funding downturn.

Emergence of statistical machine learning: Focus on learning from data rather than explicit programming.

Key algorithms: Support Vector Machines (SVMs), Decision Trees, and Bayesian networks gained prominence.

Speaker Notes

Despite the successes, the limitations of expert systems eventually led to the 'Second AI Winter.' Their brittleness and difficulty in adapting to new information became apparent. The specialized hardware supporting LISP machines also failed to gain widespread adoption, contributing to the downturn. However, beneath the surface, a new paradigm was quietly gaining traction: statistical machine learning. Researchers began to explore algorithms that could learn patterns directly from data, moving away from hand-coded rules. Algorithms like Support Vector Machines and Decision Trees proved powerful, laying the groundwork for the data-driven AI we know today.

Data, Compute, and Algorithms: The ML Explosion (2000s)

`Slide 8`

Availability of 'Big Data': Internet growth provided unprecedented amounts of digital information.

Increased computational power: Moore's Law continued, coupled with the rise of Graphical Processing Units (GPUs).

Algorithmic advancements: Ensemble methods (Random Forests, Gradient Boosting) improved predictive accuracy.

Open-source software and frameworks democratized access to ML tools.

Speaker Notes

The 2000s witnessed a perfect storm of factors that fueled the machine learning explosion. The internet generated vast quantities of data, which became the 'fuel' for learning algorithms. Simultaneously, advances in hardware, particularly the repurposing of GPUs for general-purpose computing, provided the necessary processing power. Algorithmic innovations, such as ensemble methods, further boosted performance. The rise of open-source software made these tools accessible to a wider community, leading to rapid experimentation and innovation. This era solidified the data-driven approach as the dominant paradigm in AI.

The Deep Learning Revolution (2010s)

`Slide 9`

Revival of Artificial Neural Networks (ANNs) with multiple hidden layers ('deep' networks).

Breakthroughs in training deep networks: Rectified Linear Units (ReLUs), dropout, better optimizers.

ImageNet moment (2012): AlexNet's dramatic performance improvement in image recognition.

Convolutional Neural Networks (CNNs) for image and video, Recurrent Neural Networks (RNNs) for sequential data.

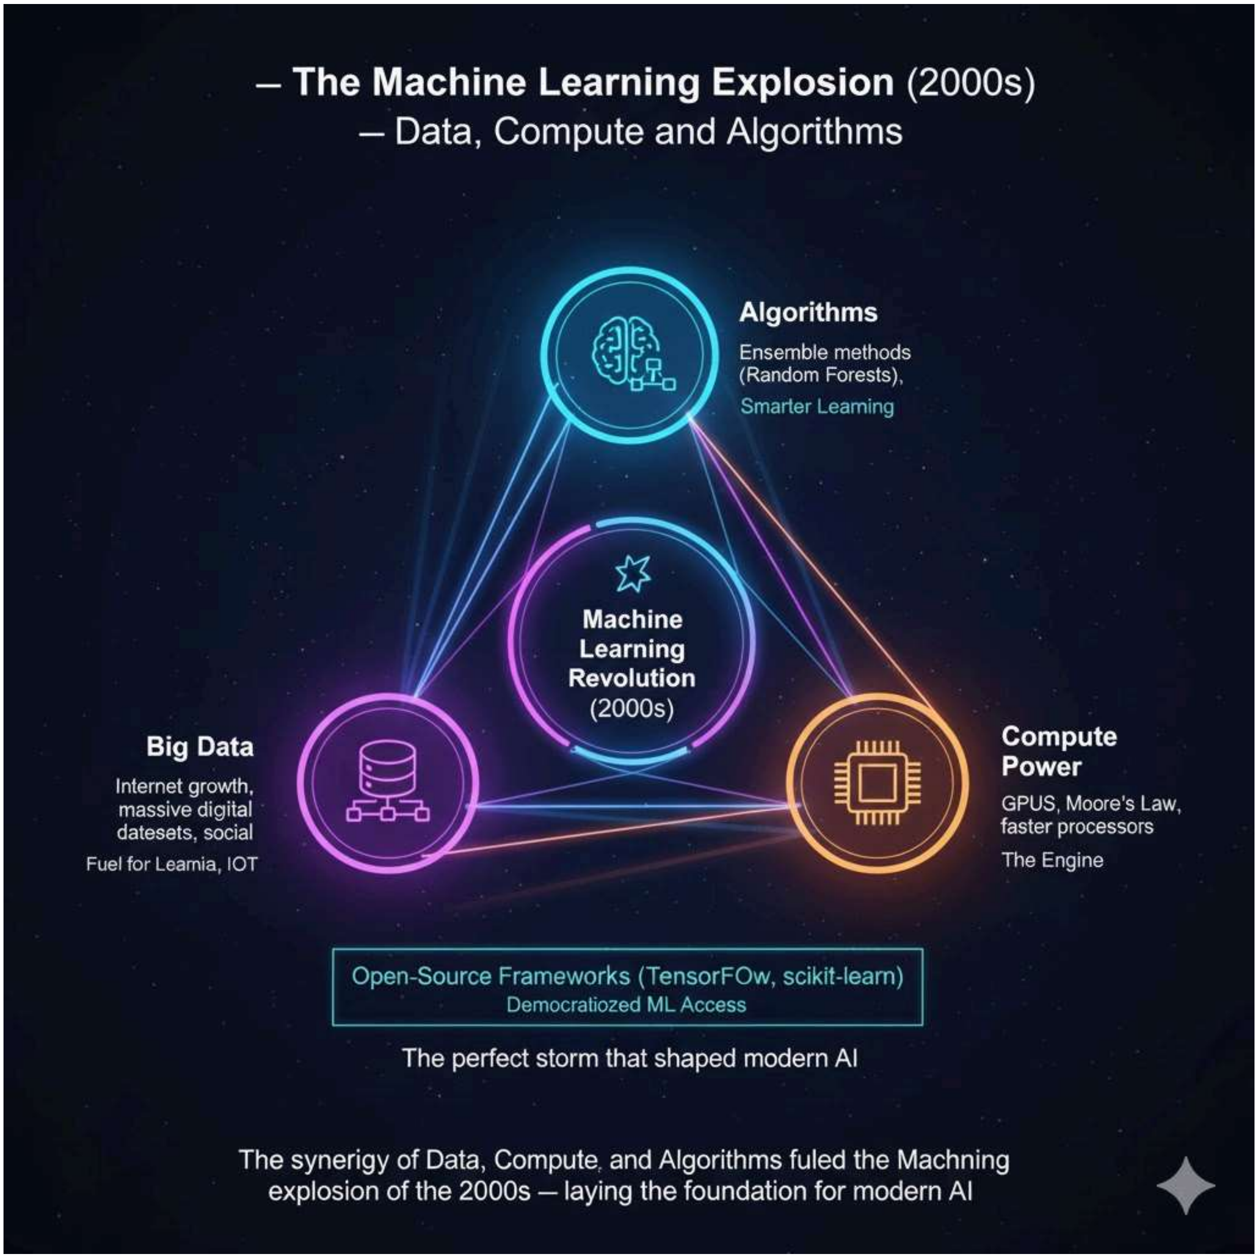

Speaker Notes

The 2010s marked the 'Deep Learning Revolution.' Artificial Neural Networks, which had seen limited success in previous decades, made a dramatic comeback. Key innovations like ReLU activation functions, dropout regularization, and improved optimization algorithms allowed for the *successful training of much deeper networks. The ImageNet Large Scale Visual Recognition Challenge in 2012, where AlexNet achieved unprecedented accuracy, was a watershed moment. Deep learning, especially CNNs for computer vision and RNNs for natural language processing, quickly became the state-of-the-art across numerous domains.*

The Transformer Architecture: A Game Changer

`Slide 10`

Introduced in 'Attention Is All You Need' (2017) by Google Brain.

Revolutionized sequence-to-sequence modeling, replacing RNNs and LSTMs.

Key innovation: Self-attention mechanism, allowing parallel processing of input sequences.

Captures long-range dependencies efficiently, crucial for natural language understanding and generation.

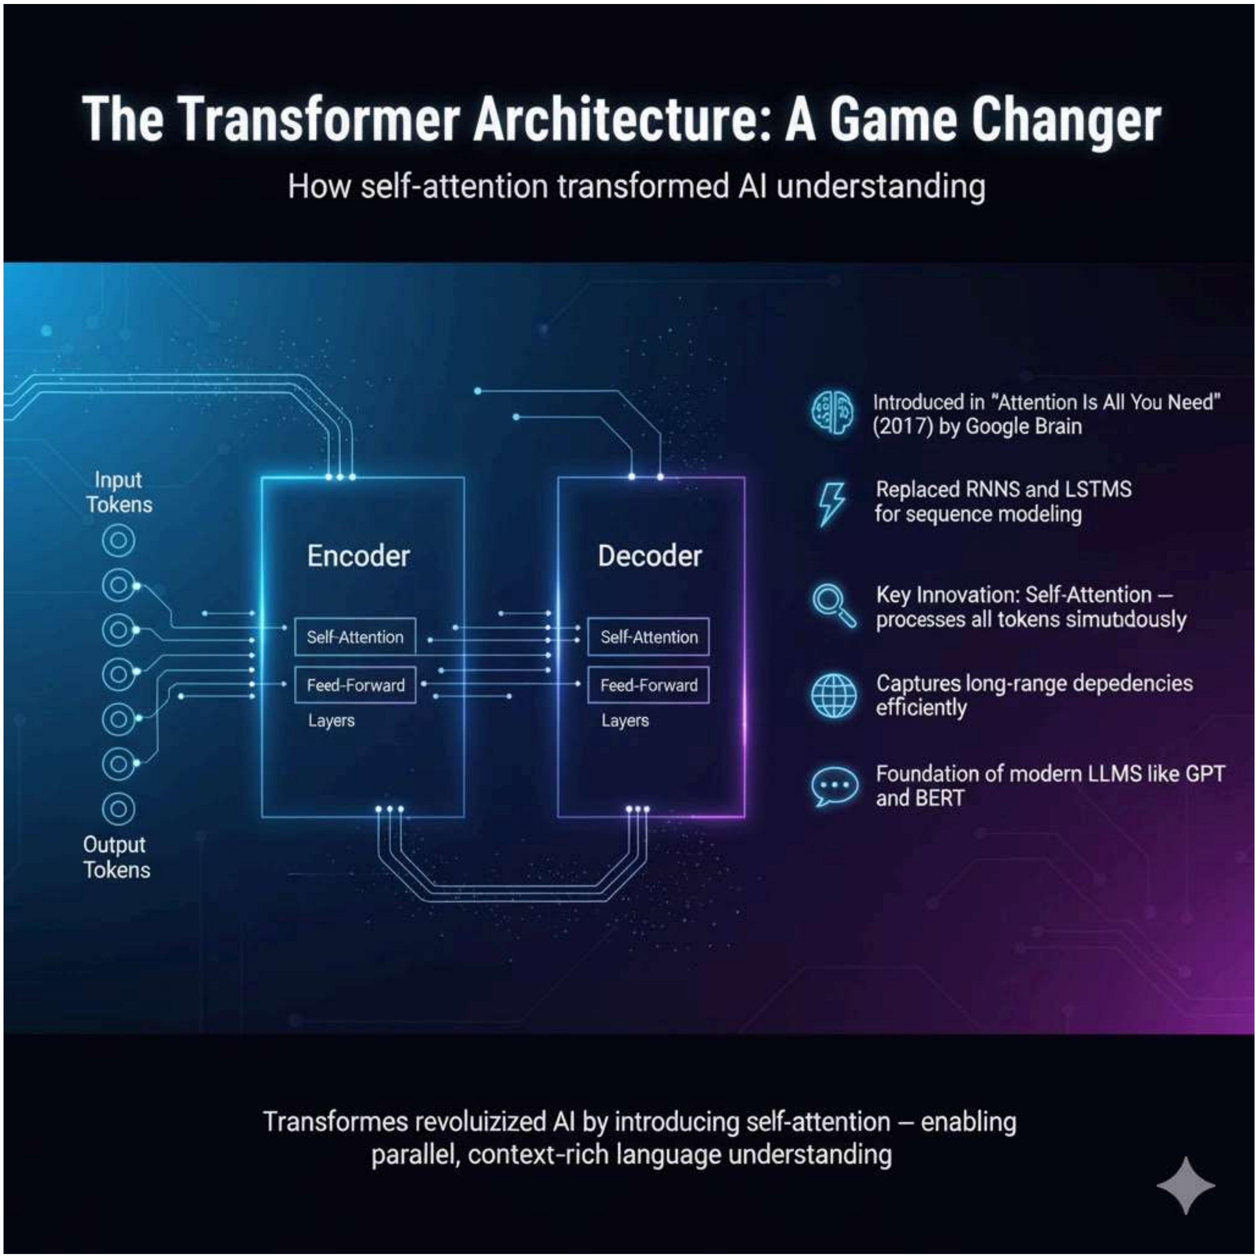

Speaker Notes

While deep learning was already transformative, the introduction of the Transformer architecture in 2017 was a true game-changer, particularly for natural language processing. It fundamentally changed how we process sequential data, moving away from the sequential nature of RNNs. Its core innovation, the self-attention mechanism, allows the model to weigh the importance of different parts of the input sequence simultaneously, enabling parallel processing. This efficiency and ability to capture long-range dependencies made it incredibly powerful for tasks like translation, text summarization, and question answering, paving the way for modern LLMs.

How Large Language Models (LLMs) Work

`Slide 11`

Trained on a vast corpora of text and code to predict the next token (word or sub-word).

Utilize the Transformer architecture to process context and generate coherent text.

Emergent abilities: With sufficient scale, LLMs exhibit capabilities not explicitly programmed, like reasoning or common sense.

Fine-tuning (e.g., instruction tuning, Reinforcement Learning from Human Feedback) refines behavior.

Speaker Notes

So, how do these powerful Large Language Models actually work? At their core, LLMs are sophisticated statistical models trained on enormous datasets of text and code. Their primary objective during training is surprisingly simple: predict the next word or token in a sequence. The Transformer architecture allows them to process vast contexts and understand complex grammatical and semantic relationships. What's truly remarkable are their 'emergent abilities' – capabilities like complex reasoning, translation, or even coding that appear only when models reach a certain scale, without explicit programming. Post-training, techniques like fine-tuning with human feedback further align their behavior with user expectations.

Key Breakthroughs Enabling Modern LLMs

`Slide 12`

Massive Datasets: Availability of web-scale text (Common Crawl, Wikipedia, books, code).

Scalable Computing Infrastructure: Advanced GPUs, TPUs, and distributed training techniques.

Transformer Architecture: Enabled parallel training and efficient handling of long sequences.

Scaling Laws: Empirical findings demonstrating predictable performance gains with increased data, model size, and computation.

Speaker Notes

The advent of modern LLMs wasn't a single breakthrough, but a confluence of several critical advancements. Firstly, the sheer volume of high-quality training data available from the internet provided the necessary fuel. Secondly, the development of specialized hardware like GPUs and TPUs, combined with sophisticated distributed training techniques, made it computationally feasible to train models with billions of parameters. The Transformer architecture, as discussed, provided the efficient model structure. Finally, the discovery of 'scaling laws' provided a roadmap, showing that with more data, larger models, and more compute, performance consistently improves in predictable ways, guiding researchers towards today's colossal models.

The Era of Agentic AI: Beyond Prediction

`Slide 13`

Moving beyond mere text generation to autonomous decision-making and action.

LLMs as 'brains' for agents, enabling reasoning, planning, and tool use.

Agentic loops: Perceive, Plan, Act, Reflect — allowing iterative improvement and task execution.

Examples: Autonomous research agents, code-generating agents, conversational assistants with memory and agency.

Speaker Notes

We are now entering the era of Agentic AI, which represents a significant paradigm shift beyond just generating text or predictions. Here, LLMs act as the 'brain' for intelligent agents, empowering them with the ability to reason, plan, and interact with tools and environments. These agents operate in *iterative 'agentic loops' – they perceive their environment, formulate a plan, execute actions, and then reflect on the outcomes to improve. This enables them to tackle complex, multi-step tasks autonomously, from conducting research to writing and debugging code, marking a step closer to truly intelligent and autonomous systems.*

Challenges and Future Directions

`Slide 14`

Ethical considerations: Bias, fairness, privacy, misuse.

Safety and alignment: Ensuring AI systems operate safely and align with human values.

Interpretability and explainability: Understanding why AI makes certain decisions.

Pursuit of Artificial General Intelligence (AGI) and superintelligence.

Integration into various industries and societal impact.

Speaker Notes

As AI continues to advance, so do the challenges. Ethical considerations around bias, fairness, and privacy are paramount. Ensuring AI systems are safe, robust, and aligned with human values is a critical area of research. Understanding the 'black box' nature of deep learning models through interpretability and explainability remains a significant hurdle. The long-term goal of Artificial General Intelligence, or AGI, which aims for human-level cognitive abilities across a wide range of tasks, continues to drive research. The integration of AI into every facet of society will undoubtedly bring profound changes and necessitate careful navigation.

Conclusion: A Journey of Continuous Innovation

`Slide 15`

history is a testament to human ingenuity and perseverance through cycles of optimism and winter.

From symbolic logic to data-dAI's riven deep learning and agentic systems, each era built upon the last.

The future promises even more transformative advancements, alongside complex ethical and societal considerations.

Continuous research and responsible development are crucial for harnessing AI's full potential.

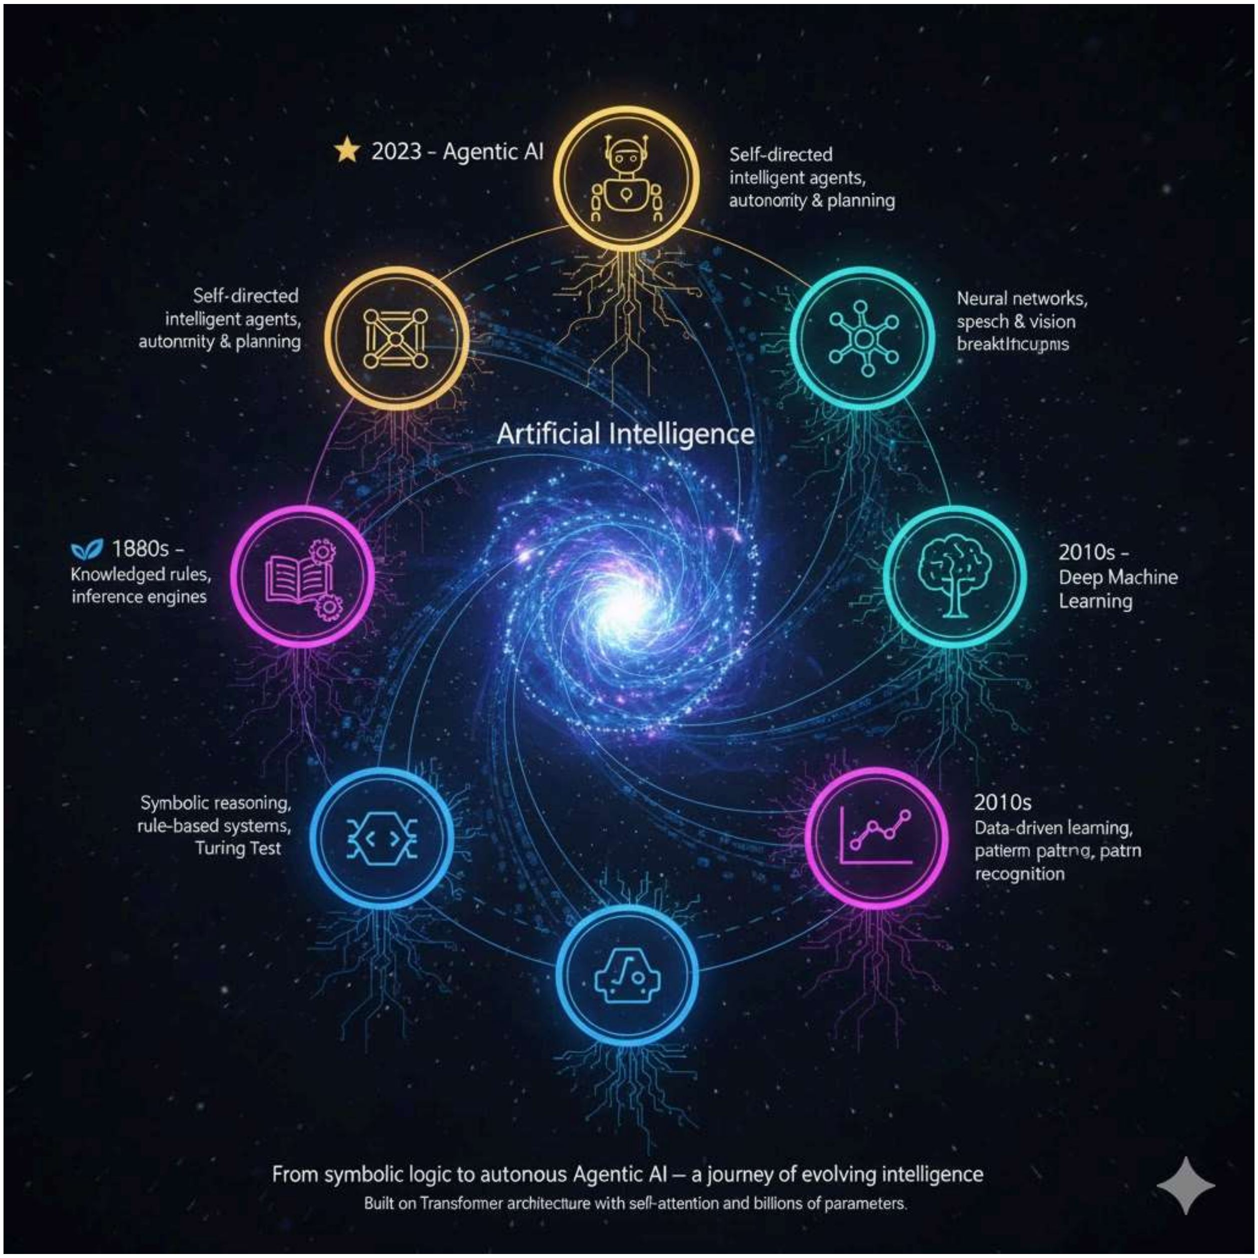

Speaker Notes

In conclusion, the history of AI is a rich tapestry woven with threads of grand ambition, scientific rigor, periods of both immense optimism and sober realism. We've seen the field evolve from abstract philosophical concepts to practical expert systems, through the data-driven revolution of machine learning and deep learning, and now into the exciting frontier of agentic AI. Each phase has contributed invaluable lessons and technologies. As we look ahead, the pace of innovation shows no signs of slowing. It is our collective responsibility to guide this powerful technology responsibly, ensuring its benefits are broadly shared and its challenges are thoughtfully addressed. Thank you.

Deep Learning Era (2010–2020) Neural networks evolved into Deep Learning — systems inspired by the human brain with multiple layers of neurons. Tip: Explain — 'Deep Learning helps AI see, hear, and think smarter.

Large Language Models (LLMs)

Models like GPT and BERT understand and generate language using billions of parameters to predict the next word in a sentence. .

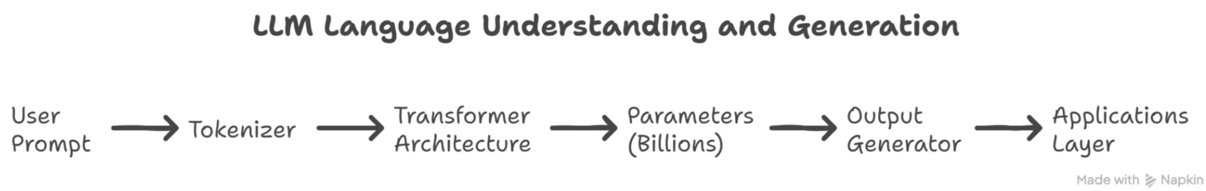

• LLMs like GPT and Gemini understand and generate human language.

• Built using the Transformer architecture with self-attention.

• Trained on billions of parameters for contextual learning. Key Breakthroughs Behind LLMs
• Transformer model (2017) revolutionized AI.

• Self-attention allowed models to focus on relevant data parts. • Parallel training enabled massive model scaling.

'LLMs are the brains behind ChatGPT'

Agentic AI ERA (2023–Now)

Agentic AI can reason, plan, use tools, and take actions — making AI more autonomous and capable of complex tasks

- *• Combines LLMs with reasoning and tool-use.*
- *• Systems can autonomously plan, act, and learn from results.*
- *• Examples: GPTs, AutoGPT, LangChain agents*

AI Timeline Overview

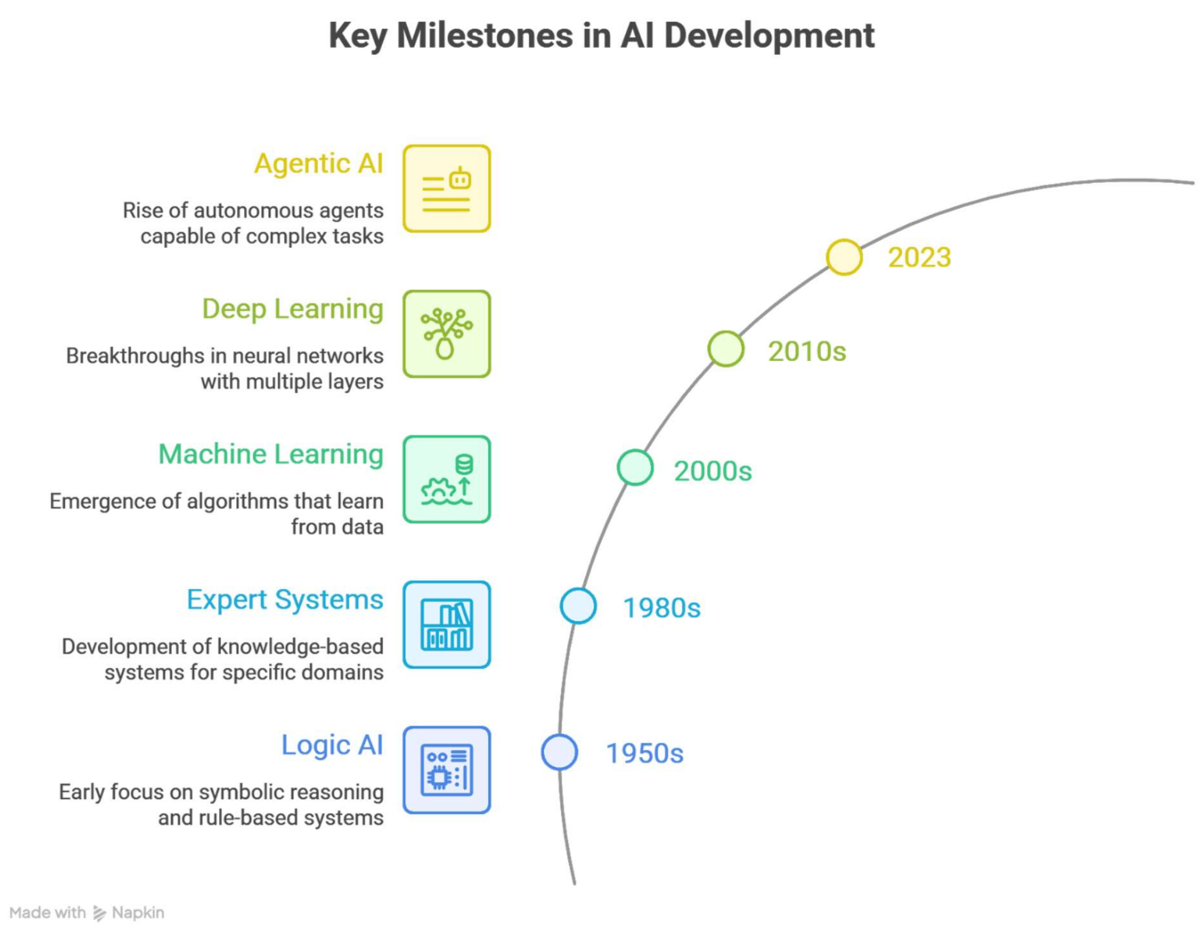

Timeline: 1950s (Logic AI) → 1980s (Expert Systems) → 2000s (ML) → 2010s (Deep Learning) → 2023 (Agentic AI

Deep Learning Architecture

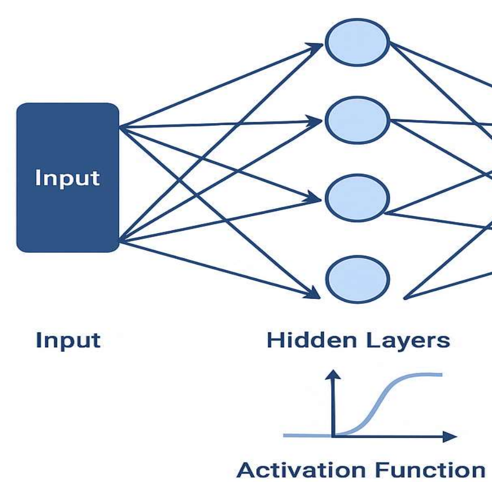

Conclusion

From logic-based systems to Agentic AI, the journey of Artificial Intelligence reflects human creativity and innovation.

'AI is no longer just a dream; it’s our digital partner’.

Prepared by: Vandana Chohan
Roll No: 00304612
Assignment 1
Quarter 4
Topic:Evolution of Artificial Intelligence.